

Other agents prove they can trade. Seoul Shield proves which safety policy deserves the right to trade.

\$100,000 PAPER CASH / EQUITY	245 SPY OPTION CONTRACTS	21 / 80 / 144 OPPORTUNITY / UNCERTAIN / BLOCKED	\$1,096 MAX LOSS
---	------------------------------------	---	----------------------------

THE PROBLEM

AI can produce a persuasive options thesis, but persuasion is not permission to move capital. Seoul Shield tests one signal against four safety policies and reveals why a proposal advances, resizes, or stops.

ONE SNAPSHOT, FOUR POLICIES

One fixed, hash-identified snapshot feeds every policy and links to the audit chain. No Guard preserves the proposal as a shadow baseline. Static Guard applies fixed limits. Adaptive Guard evaluates regime, quote age, liquidity, exposure, and recent loss. Live Execution is the only path with Paper authority, still gated by deterministic rules and human approval.

HARD SAFETY INVARIANTS

- Two-leg, same-expiration debit spreads only
- 1% maximum trade loss; 5% portfolio risk; 2% daily stop
- Strike ordering and position_intent validation
- Unique client_order_id and duplicate prevention
- Model output can never relax a hard limit
- Hash-chain audit of every decision and preview

VERIFIED ALPACA REPLAY

An active Paper account showed \$100,000 cash/equity, \$400,000 buying power, Options Level 3, and zero positions or orders. The engine inspected 245 SPY calls expiring September 4, 2026 and sorted them into 21 Opportunity, 80 Uncertain, and 144 Risk Blocked.

THE DECISION

Candidate gates converged on BUY SPY 765C at the \$5.06 ask and SELL SPY 770C at the \$2.32 bid. Four contracts at a \$2.74 net debit create \$1,096 maximum loss. The deterministic limit is \$1,000, so the capital gate closes at \$96 over limit.

HARD GATE: BLOCKED

QUOTE STALE | FALLBACK / NOT LIVE AI
PAPER PREVIEW | NOT SUBMITTED

No fill, return, or realized P&L; is claimed. The public app is a sanitized read-only evidence room, not an execution console.

ALPACA	IMPLEMENTATION EVIDENCE	PUBLIC MODE
Trading + Market Data APIs Read-only MCP evidence Paper MLEG payload	Python risk core: 20 tests Idempotency + status/cancel paths Three.js tribunal + 2D fallback	VERIFIED ALPACA REPLAY PAPER PREVIEW ALPACA_ALLOW_SUBMIT=false

245 → 21 / 80 / 144 → 21 → 8 → 4 → 2 → **BLOCKED**

 LIVE DEMO https://seoul-shield-risk-tribunal.ljs2546.chatgpt.site	 GITHUB https://github.com/lsh2546/seoul-shield-counterfactual-risk-tribunal	 YOUTUBE https://www.youtube.com/watch?v=nveiZhnRdK8
---	--	---

Educational hackathon prototype - not investment advice - no live-money account support